

```
In [1]: import pandas as pd
```

```
In [2]: data=pd.read_csv("/home/placement/Desktop/reddy/flat500.csv")
```

```
#data['model']=data['model'].map({'lounge':1,'pop':2,'sport':3})
```

```
In [3]: data1=data.drop(['ID','lat','lon'],axis=1)
```

```
In [4]: data1
```

```
Out[4]:
```

	model	engine_power	age_in_days	km	previous_owners	price
0	lounge	51	882	25000	1	8900
1	pop	51	1186	32500	1	8800
2	sport	74	4658	142228	1	4200
3	lounge	51	2739	160000	1	6000
4	pop	73	3074	106880	1	5700
...	...	...	...	...	...	...
1533	sport	51	3712	115280	1	5200
1534	lounge	74	3835	112000	1	4600
1535	pop	51	2223	60457	1	7500
1536	lounge	51	2557	80750	1	5990
1537	pop	51	1766	54276	1	7900

1538 rows × 6 columns

Type *Markdown* and LaTeX:  $\alpha^2$

```
In [5]: #data1=pd.get_dummies(data1)
```

```
In [6]: #data1
```

```
In [7]: data1['model']=data1['model'].map({'lounges':1,'pop':2,'sport':3})
```

```
In [8]: data1
```

```
Out[8]:
```

	model	engine_power	age_in_days	km	previous_owners	price
0	1	51	882	25000	1	8900
1	2	51	1186	32500	1	8800
2	3	74	4658	142228	1	4200
3	1	51	2739	160000	1	6000
4	2	73	3074	106880	1	5700
...	...	...	...	...	...	...
1533	3	51	3712	115280	1	5200
1534	1	74	3835	112000	1	4600
1535	2	51	2223	60457	1	7500
1536	1	51	2557	80750	1	5990
1537	2	51	1766	54276	1	7900

1538 rows × 6 columns

```
In [9]: data1=pd.get_dummies(data1)
```

```
In [10]: data1
```

```
Out[10]:
```

	model	engine_power	age_in_days	km	previous_owners	price
0	1	51	882	25000	1	8900
1	2	51	1186	32500	1	8800
2	3	74	4658	142228	1	4200
3	1	51	2739	160000	1	6000
4	2	73	3074	106880	1	5700
...	...	...	...	...	...	...
1533	3	51	3712	115280	1	5200
1534	1	74	3835	112000	1	4600
1535	2	51	2223	60457	1	7500
1536	1	51	2557	80750	1	5990
1537	2	51	1766	54276	1	7900

1538 rows × 6 columns

```
In [11]: y=data1['price']  
x=data1.drop('price',axis=1)
```

In [12]:

y

Out[12]:

```
0      8900
1      8800
2      4200
3      6000
4      5700
```

```
...
1533   5200
1534   4600
1535   7500
1536   5990
1537   7900
```

Name: price, Length: 1538, dtype: int64

In [13]:

x

Out[13]:

	model	engine_power	age_in_days	km	previous_owners
0	1	51	882	25000	1
1	2	51	1186	32500	1
2	3	74	4658	142228	1
3	1	51	2739	160000	1
4	2	73	3074	106880	1
...	...	...	...	...	...
1533	3	51	3712	115280	1
1534	1	74	3835	112000	1
1535	2	51	2223	60457	1
1536	1	51	2557	80750	1
1537	2	51	1766	54276	1

1538 rows × 5 columns

```
In [14]: #!/pip3 install scikit-learn
```

```
In [15]: #data1=pd.get_dummies(data1)
```

```
In [16]: #data1
```

```
In [ ]:
```

```
In [17]: from sklearn.model_selection import train_test_split  
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.33,random_state=42)
```

```
In [18]: x_test.head(5)
```

Out[18]:

	model	engine_power	age_in_days	km	previous_owners
481	2	51	3197	120000	2
76	2	62	2101	103000	1
1502	1	51	670	32473	1
669	1	51	913	29000	1
1409	1	51	762	18800	1

```
In [19]: x_train.shape
```

Out[19]: (1030, 5)

```
In [20]: y_train.shape
```

Out[20]: (1030,)

```
In [21]: x_train.head()
```

```
Out[21]:
```

	model	engine_power	age_in_days	km	previous_owners
527	1	51	425	13111	1
129	1	51	1127	21400	1
602	2	51	2039	57039	1
331	1	51	1155	40700	1
323	1	51	425	16783	1

```
In [22]: y_train.head()
```

```
Out[22]: 527    9990
129     9500
602    7590
331    8750
323    9100
Name: price, dtype: int64
```

```
In [23]: x_test.head()
```

```
Out[23]:
```

	model	engine_power	age_in_days	km	previous_owners
481	2	51	3197	120000	2
76	2	62	2101	103000	1
1502	1	51	670	32473	1
669	1	51	913	29000	1
1409	1	51	762	18800	1

In [24]: `y_test.head()`

Out[24]:

481	7900
76	7900
1502	9400
669	8500
1409	9700

Name: price, dtype: int64

In [25]: `from sklearn.linear_model import LinearRegression  
reg=LinearRegression()  
reg.fit(x_train,y_train)`

Out[25]:

▼ LinearRegression

LinearRegression()

In [26]: `#!pip3 install scikit-learn`

In [27]: `ypred=reg.predict(x_test)`

In [28]: ypred

```
Out[28]: array([ 5994.51703157,  7263.58726658,  9841.90754881,  9699.31627673,
                10014.19892635,  9630.58715835,  9649.4499026 , 10092.9819664 ,
                9879.19498711,  9329.19347948, 10407.2964056 ,  7716.91706011,
                7682.89152522,  6673.95810983,  9639.42618839, 10346.53679153,
                9366.53363673,  7707.90063494,  4727.33552438, 10428.17092937,
                10359.87663878, 10364.84674179,  7680.16157493,  9927.58506055,
                7127.7284177 ,  9097.51161986,  4929.31229715,  6940.60225317,
                7794.35120591,  9600.43942019,  7319.85877519,  5224.05298205,
                5559.52039134,  5201.35403287,  8960.11762682,  5659.72968338,
                9915.79926869,  8255.93615893,  6270.40332834,  8556.73835062,
                9749.72882426,  6873.76758364,  8951.72659758, 10301.95669828,
                8674.89268564, 10301.93257222,  9165.73586068,  8846.92420399,
                7044.68964545,  9052.4031418 ,  9390.75738772, 10267.3912561 ,
                10046.90924744,  6855.71260655,  9761.93338967,  9450.05744337,
                9274.98388541, 10416.00474283,  9771.10646661,  7302.96566423,
                10082.61483093,  6996.96553454,  9829.40534825,  7134.21944391,
                6407.26222178,  9971.82132188,  9757.01618446,  8614.84049875,
                8437.92452169,  6489.24658616,  7752.65456507,  6626.60510856,
                8329.88998217, 10412.00324329,  7342.77348105,  8543.63624413,
                8706.44742777, 10010.42502651,  7256.86706062,  8522.1400051 ])
```

In [29]: `from sklearn.metrics import r2_score`  
`r2_score(y_test,ypred)`

Out[29]: 0.8383895235218546

In [30]: `from sklearn.metrics import mean_squared_error as ms`  
`o=ms(y_test,ypred)`  
`o`

Out[30]: 593504.2888137395

In [31]: `import math`  
`math.sqrt(o)`

Out[31]: 770.3922954013361



In [32]: ypred

```
Out[32]: array([ 5994.51703157,  7263.58726658,  9841.90754881,  9699.31627673,
 10014.19892635,  9630.58715835,  9649.4499026 , 10092.9819664 ,
   9879.19498711,  9329.19347948, 10407.2964056 ,  7716.91706011,
   7682.89152522,  6673.95810983,  9639.42618839, 10346.53679153,
   9366.53363673,  7707.90063494,  4727.33552438, 10428.17092937,
 10359.87663878, 10364.84674179,  7680.16157493,  9927.58506055,
   7127.7284177 ,  9097.51161986,  4929.31229715,  6940.60225317,
   7794.35120591,  9600.43942019,  7319.85877519,  5224.05298205,
   5559.52039134,  5201.35403287,  8960.11762682,  5659.72968338,
   9915.79926869,  8255.93615893,  6270.40332834,  8556.73835062,
   9749.72882426,  6873.76758364,  8951.72659758, 10301.95669828,
   8674.89268564, 10301.93257222,  9165.73586068,  8846.92420399,
   7044.68964545,  9052.4031418 ,  9390.75738772, 10267.3912561 ,
 10046.90924744,  6855.71260655,  9761.93338967,  9450.05744337,
   9274.98388541, 10416.00474283,  9771.10646661,  7302.96566423,
 10082.61483093,  6996.96553454,  9829.40534825,  7134.21944391,
   6407.26222178,  9971.82132188,  9757.01618446,  8614.84049875,
   8437.92452169,  6489.24658616,  7752.65456507,  6626.60510856,
   8329.88998217, 10412.00324329,  7342.77348105,  8543.63624413,
   8786.44742777, 10010.42502651,  7256.86706062,  8522.14000511])
```

```
In [33]: Results=pd.DataFrame(columns=['price','predicted'])
Results['price']=y_test
Results['predicted']=ypred
Results=Results.reset_index()
Results['ID']=Results.index
Results.head(15)
```

Out[33]:

	index	price	predicted	ID
0	481	7900	5994.517032	0
1	76	7900	7263.587267	1
2	1502	9400	9841.907549	2
3	669	8500	9699.316277	3
4	1409	9700	10014.198926	4
5	1414	9900	9630.587158	5
6	1089	9900	9649.449903	6
7	1507	9950	10092.981966	7
8	970	10700	9879.194987	8
9	1198	8999	9329.193479	9
10	1088	9890	10407.296406	10
11	576	7990	7716.917060	11
12	965	7380	7682.891525	12
13	1488	6800	6673.958110	13
14	1432	8900	9639.426188	14

```
In [34]: Results['price_diff']=Results.apply(lambda row: row.price - row.predicted,axis=1)
```

In [35]: Results

Out[35]:

	index	price	predicted	ID	price_diff
0	481	7900	5994.517032	0	1905.482968
1	76	7900	7263.587267	1	636.412733
2	1502	9400	9841.907549	2	-441.907549
3	669	8500	9699.316277	3	-1199.316277
4	1409	9700	10014.198926	4	-314.198926
...	...	...	...	...	...
503	291	10900	10007.364639	503	892.635361
504	596	5699	6390.174715	504	-691.174715
505	1489	9500	10079.478928	505	-579.478928
506	1436	6990	8363.337585	506	-1373.337585
507	575	10900	10344.486077	507	555.513923

508 rows × 5 columns

In [ ]: